

Application of Azure Blob Storage and Apache Spark in Economic Indicators Trend Analysis

Manuel Benitez
MSc. In Data Analytics
National College of Ireland
Dublin, Ireland

Abstract – A scalable pipeline has been developed which integrates Azure Blob Storage for data extraction, PySpark for distributed processing, Cosmos DB for analytical storage and Power BI for visualisation. The pipeline handles five datasets: Gross Domestic Product (GDP), labour force, employment, population and gender statistics. Including data from 1960 to 2024 inclusive across 264 countries and regions. Schema implementation, temporal alignment and hierarchical reshaping was implemented for analytical purposes. The system demonstrates how cloud object storage combined with distributed processing frameworks can overcome traditional limitations in economic data analysis. Power BI makes possible to compare economies within different regions and timeframes. This research contributes a replicable pattern for multidimensional economic analysis while addressing key challenges in data versioning and schema evolution.

Keywords – Big Data, Azure Blob Storage, PySpark, Azure Cosmos DB,

1. INTRODUCTION

The global economy is characterized by its complexity and multidimensionality trends that require robust and scalable analytics to uncover actionable insights. Traditional economic analysis is often hampered by data silos, inconsistent schemas issues, and limited computational resources. This project aims to address these challenges by constructing a cloud-based data pipeline capable of large-scale economic and demographic datasets ingestion, processing, and visualisation .

The central research question is:

“How can cloud-native big data technologies be utilized to support scalable, multidimensional analysis of macroeconomic and gender trends in different economies?”

The research aims to combine World Bank economic indicators in Spark-optimized Azure architecture. Understanding macroeconomic and gender-related indicators

at scale is critical for informing policy, investment, and development strategies. By integrating diverse datasets and enabling interactive exploration, this research supports data-driven decision-making for governments, NGOs, and researchers.

And the motivation relies on the need of understanding how different macroeconomic indicators.

2. RELATED WORK

2.1 Big Data

Big data plays an essential role in the analysis of economics due to its ability to transform vast, diverse datasets into actionable insights, but this potential represents significant challenges. Traditional analytical methods struggle with the volume, velocity, and variety of economic data, leading to issues in storage, processing, and integration. While more businesses are storing big data, they are yet to utilize them to their full potential and are therefore yet to realize their full impacts [1]. Some solutions to this problem are developed by organizations, to adopt advanced technologies such as distributed computing frameworks, cloud storage, and real-time analytics, such as robust infrastructures using cloud services and leverage tools like Hadoop and Spark for scalable, efficient data processing. Investing in in data cleansing, integration, and governance to ensure data quality and consistency, while it requires of security protocols implementation to protect sensitive information. These strategies empower both businesses and governments to improve big data for decision-making, risk management, and operational efficiency, ultimately supporting economic growth and innovation. [2].

2.2 Distributed Storage Strategy and Visual Analysis for Economic Big Data

Cloud technologies transform data storage and processing capabilities (i.e. the use of Azure cloud services). This transformation is evident in the growth of the cloud sector and increasing market concentration, highlighting the rising reliance on Azure platform from Microsoft. The ability of computers systems to exchange and use information and

network effects in cloud ecosystem are relevant to the integration of different components. Therefore, a specialised storage architecture for economic big data was created based on its sources and characteristics, analysing economic data types and proposed a unified coding approach based on spatiotemporal data multilevel division strategy. Using MongoDB to verify the performance of their multilevel partition algorithm. The system was designed to address challenges with large numbers of small files and diverse file types in economic big data storage and retrieval that improved storage performance, while reducing storage management costs, with no storage space limitation [3].

2.3 Azure Blob Storage

Azure Blob Storage remains as one of the most widely known solutions for cloud storage providers specifically for hybrid cloud environment, however it shows variability in performance depending on the selected storage tier integrated with Microsoft Defender for Cloud which makes it capable of strong security insights [4]. Moreover, it provides advanced data cycle management features, which allows automation data retention and archiving strategies. It also provides deep integration with Azure Synapse Analytics, Power BI and Azure Machine Learning [4].

Azure Blob Storage is compatible with many popular development frameworks, including Python, Java, and .NET. It stands out as the only cloud storage service offering premium solid-state drive (SSD)-based object storage, designed for low-latency and interactive workloads.

2.4 Application of Big Data and Cloud Computing in Economic Management Database

Big data and cloud computing technologies can improve the construction of economic database for a more robust management. The research highlights the need for big data dissemination and its positive effect on improving the effectiveness and efficiency of financial management, as a support to enable data-driven decision-making. The development includes the analysis of economic big data types based on source and characteristics, the design of database storage architecture and data storage structure, to implement a unified coding method for spatiotemporal data and development of a multilevel partition algorithm. MongoDB was used as tool for the construction of a prototype system. Followed by the implementation of Wiener distributed memory for workload storage and the division of workloads into distributed storage windows of specific duration. The main challenges were the difficulty in implementing distributed storage to meet all requirements simultaneously, and the large numbers of small files and diverse file types in economic data [5].

2.5 The Economics of the Cloud

The use cloud-based infrastructure to handle large datasets with varying structures by improving a multi-layered architecture that separates storage, processing, and analysis functions. Distributed computing frameworks were employed to process economic data at scale. It demonstrated how cloud technologies enable cost-effective automation of data processing workflows, essential for pipelines that process multiple economic datasets simultaneously [6].

3. METHODOLOGY

The CRISP-DM methodology [7] phases applied are:

Business Understanding. Defining project objectives and requirements.

Data Understanding. Accessing the data to describe, explore, and verify its quality.

Data Preparation. Adequately prepare the data for loading stage by selecting, cleaning, formatting the data, and feature engineering to find new insights within the data.

Modelling. Selecting and applying appropriate such predictive or descriptive models.

Evaluation (Results). Assessing the data processing skills of the pipeline to meet research objectives.

Deployment (Results). Deploying the final model into a real-world database and enable it for visualisation.

This research evaluated applicable CRISP-DM [8] stages, however, some names were changes and modelling hasn't been applied due to the nature of the case of study.

3.1 Business Understanding

Analyse macroeconomic and gender trends across countries over time. The central research question is:

“How can cloud-native big data technologies be utilized to support scalable analysis of macroeconomic and gender trends in different economies?”

Which raises the secondary research questions about the data from the World Bank and the global economic and gender-related indicators:

- What are the main challenges in harmonisation and temporally alignment homogeneous economic and demographic datasets?
- How do macroeconomic indicators such as GDP, employment, labour force, population and gender evolve over time from 1995 to 2024.
- What are the labour disparities in gender-related economic indicators (such as labour force,

unemployment, employment and education attainment rates) across regions over time.

3.2 Data Understanding

Five datasets were sourced from the World Bank repository, each providing annual observations from 1960 to 2024 for up to 266 countries and regions:

Dataset	Records	Dimensions	Temporal Scope
GDP	266	266 countries × 63 years	1960-2023
Labor	266	266 economies × 50 years	1974-2024
JOBS	40,177	243 countries x 167 series of indicators	1990-2016
Population	40,703	263 countries × 154 series of indicators	1960-2023
Gender	306,345	266 countries x 1159 series of indicators	1990-2023

Table 1. Datasets description.

Each dataset included a wide range of indicators, with columns for country, indicator, and annual values. The data was highly homogeneous, with varying schemas, missing values, and inconsistent year coverage.

3.3 Data Preparation

a. Pipeline Architecture

Based on a classic Extract-Transform-Load (ETL) pattern, the pipeline was implemented is shown in Fig 1.




Figure 1. Processing Pipeline of the research.

b. Extraction

Different data files from the World Bank were stored in Azure Blob Storage and accessed via secure SAS token. PySpark was used to read each file, including inferring initial schema, and standardize all the year columns involved.

c. Transformation

The key steps involved in the transformation process are as follows:

- **Schema reading.** Each CSV file was read with headers from the Azure Blob Storage through a Spark Session using Azure credentials.
- **Schema harmonization.** All datasets were standardized to have a common schema design [9] [10], the new columns names are listed in Table 2, to ensure data quality across datasets.

Old Schema	New schema	Output Datatype
Country Name	CountryName	String
Country Code	CountryCode	String
Series/Indicator Name	SeriesName	String
Series Code	SeriesCode	String
--	SeriesCategory	String
Years 1960–2024, i.e. • 1990 • 2024 [2024]	1960 – 2024	Float

Table 2. Schema design and datatypes for all datasets

- **Temporal alignment.** Not all datasets contain data from the same number of years, these were added with default values (0), ensuring consistent time series across datasets.
- **Hierarchical reshaping.** The capacity of World Bank is limited for the number of records to download, and Gender datasets were originally split across multiple files from different yeas, these were merged and restructured for unified analysis.
- **Data cleaning.** Null and missing values were filled, and non-numeric entries in year columns were coerced to floats. Also, for research purposes, the data was delimited from 1995 to 2024 as last thirty years, due to the lack of data previous that period across different datasets. The output dataset contains 387742 records and 35 variables.
- **ID creation.** Unique IDs were created by concatenating country and indicator codes, as part of

the requirements to upload to Azure Cosmos DB as a dictionary.

d. Loading

Processed data was loaded into Azure Cosmos DB using the Python Software Development Kit. Each record was inserted with robust error handling and retry logic, which means if an error occurred (i.e., due to a network failure), the operation would retry up to a maximum of three times by default, while failures were tracked for further inspection.

e. Visualisation

Power BI was connected to Cosmos DB using the Azure Cosmos DB connector in import mode [11] [12]. The analytical store schema enabled direct querying and dashboarding without impacting transactional workloads, pivoting was essential to produce time series and visualise trends within the indicators and countries across the series of study.

Year columns were unpivoted to create a long-form table (Year, Value), enabling time series analysis in Power BI, Providing more than 11millions of rows for visualisation and analysis purposes.

f. Automation

The entire workflow was automated via Python scripts, with modular functions for each ETL step. Progress bars and logging provided transparency during long-running uploads.

4. RESULTS

The approach outlined here provides a blueprint for large-scale, cloud-native economic analytics, addressing key challenges in data integration, versioning, and schema management.

Regarding its performance impact, running analytical queries directly against Cosmos DB can impact transactional workloads and consume significant request units (RUs).

Schema Changes: Changes in data source schemas can break pipelines and require manual intervention.

Cost Management: Inefficient pipelines can lead to unexpected costs, especially when processing large volumes of data.

Latency in Data Synchronization: Delays in syncing data between sources and destinations can lead to outdated insights.

Missing Data Handling: Addressing gaps in historical economic data, especially for developing economies.

Homogeneous Data Integration: Merging datasets with different structures, time periods, and granularities.

Version Control and Change Management: Managing changes to pipeline logic and tracking versions of configurations.

The final unified dataset contained:

- **387,742 rows**
- **83,049 distinct indicators**
- **Coverage:** 1995–2024, 264 countries/regions, 5 major domains (GDP, jobs, population, labour, gender)

Power BI Analysis

1. Macroeconomic Trends

- **GDP Growth:** Interactive line charts enabled comparison of GDP trajectories across regions and decades, revealing periods of rapid growth and stagnation.
- **Labour Force Participation:** Bar charts and heatmaps highlighted shifts in labour force participation, with notable increases in emerging economies post-2000.

2. Population Dynamics

- **Population Growth:** Area charts visualized demographic expansion, with filters for region and income group.
- **Dependency Ratios:** Calculated fields enabled analysis of youth and elderly dependency ratios, supporting policy analysis.

3. Gender Disparities

- **Education and Employment:** Slicers for gender and indicator allowed rapid comparison of male/female attainment in education and workforce participation.
- **Temporal Gaps:** Line charts illustrated narrowing (or persistent) gender gaps in literacy and employment over time.

4. Custom Queries


- **Country Drilldown:** Users could drill down to individual countries to view their complete economic and demographic history.
- **Regional Aggregation:** Aggregated views enabled comparison of regions (e.g., Sub-Saharan Africa vs. East Asia).

Key Findings


- **Divergent Growth Paths:** While some developing economies exhibited sustained GDP growth, others experienced volatility linked to political and external shocks.
- **Labour Force Shifts:** Female labour force participation increased in many regions, though gaps remain.
- **Population Pressure:** Rapid population growth in Africa and South Asia presents both opportunities and challenges for economic planning.
- **Data Quality:** The harmonization process revealed significant gaps and inconsistencies in international datasets, underscoring the need for robust preprocessing.
- **Data Gaps:** Deploy IoT sensors or mobile surveys to fill missing data (e.g., rural women's unpaid labour metrics).

By harnessing these technologies, stakeholders can uncover actionable insights to drive equitable economic policies and track progress toward SDGs like gender equality (SDG 5) and reduced inequalities (SDG 10).


Top 5 Countries Annual GDP growth (%) by Year




Central America Annual GDP growth (%) by Year




Countries with Highest Total Male Population (%) by Year




Countries with Highest Total Female Population (%) by Year



North America Total Labor Force by Year





4. CONCLUSIONS AND FUTURE WORK

The pipeline demonstrates that cloud-native architectures can democratize access to large-scale economic analytics. By automating schema harmonization and temporal alignment, researchers and policymakers can focus on insight generation rather than data wrangling.

Among the limitations of this project, data Gaps represent a challenge to harmonize, some countries and years remain underreported. Also, no Machine Learning as the analysis is limited to descriptive and comparative statistics; predictive modelling could be explored in future work to understand patterns and generate insights about trends in coming years to develop policies to address any gender-related indicators. Finally, the performance was a key limitation as uploading large datasets to Cosmos DB, while robust, is time-consuming and in the future, it may incur costs at scale.

Unlike traditional ETL pipelines, this approach leverages distributed processing for scalability and uses NoSQL storage to accommodate the schema evolution and homogeneous data but express the same result as approaches using MongoDB in speed. The integration with Power BI provides a user-friendly interface for non-technical stakeholders.

This project delivers a replicable, scalable pipeline for economic data analysis, integrating cloud storage, distributed processing, NoSQL analytics, and interactive dashboarding. The system enables exploration of macroeconomic and gender trends, supporting the research questions.

Future work could include the incorporation of machine learning models for forecasting, anomaly detection, and correlation between different dimensions. Also, the ingestion of real-time data using a streaming architecture to provide real-time indicators data to generate a quicker response to policy makers. Moreover, the integration with additional data sources such as health, education, climate could provide an overall panorama of the situation of different economies around the world.

References

- [geeksforgeeks.org, “Overcoming Big Data
1 Challenges: Strategies for Efficient Management
] and Analysis,” 2025. [Online]. Available:
[https://www.geeksforgeeks.org/overcoming-big-
data-challenges-strategies-for-efficient-
management-and-analysis/](https://www.geeksforgeeks.org/overcoming-big-data-challenges-strategies-for-efficient-management-and-analysis/). [Accessed 1 4
2025].
- [D. M. SalahEldin Fahmy, “The Value of Big Data
2 to the World Economy,” *The Value of Big Data to
] the World Economy*, vol. 40, no. 4, pp. 307-322,
2020.
- [X. Chang and H. Cui, “Distributed Storage
3 Strategy and Visual Analysis for Economic Big
] Data,” *Journal of Mathematics*, vol. 2021, pp. 1-
13, 2021.
- [G. Noah, “Comparing Cloud Storage Solutions
4 for Big Data: Amazon S3 vs. Google Cloud
] Storage vs. Azure Blob Storage,” 2024.
- [X. Jin, “Application of Big Data and Cloud
5 Computing in the Construction of Economic
] Management Database,” *Public Management
and Intelligent Society*, vol. 8, pp. 898-905, 2023.
- [G. Biglaiser, J. Crémer and A. Mantovani, “The
6 Economics of the Cloud,” *Toulouse School of
] Economics, Toulouse*, 2024.
- [S. Chumbar, “The CRISP-DM Process: A
7 Comprehensive Guide,” 22 09 2023. [Online].
] Available:
[https://medium.com/@shawn.chumbar/the-
crisp-dm-process-a-comprehensive-guide-
4d893aecb151](https://medium.com/@shawn.chumbar/the-crisp-dm-process-a-comprehensive-guide-4d893aecb151). [Accessed 11 12 2024].

- [Society of Actuaries in Ireland, “CRISP-DM - a
8 structured approach to planning a data
] analytics project.,” Society of Actuaries in
Ireland, Carlow, 2021.

- [N. Sayyed, “Data harmonization: Steps and best
9 practices,” 7 9 2023. [Online]. Available:
] <https://datavid.com/blog/data-harmonization>.
[Accessed 20 3 2025].

- [ApacheSpark,
1 “pyspark.sql.DataFrame.schema,” [Online].
0 Available:
] <https://spark.apache.org/docs/latest/api/python/reference/pyspark.sql/api/pyspark.sql.DataFrame.schema.html>. [Accessed 20 3 2025].

- [“Visualize Azure Cosmos DB data using Power
1 BI,” 24 9 2024. [Online]. Available:
1 [https://learn.microsoft.com/en-](https://learn.microsoft.com/en-us/azure/cosmos-db/nosql/powerbi-visualize)
] [us/azure/cosmos-db/nosql/powerbi-visualize](https://learn.microsoft.com/en-us/azure/cosmos-db/nosql/powerbi-visualize).
[Accessed 2 4 2025].

- [R. Souza, “Analytics and Business Intelligence
1 (BI) on your Azure Cosmos DB data,” 11 3 2025.
2 [Online]. Available:
] [https://github.com/MicrosoftDocs/azure-](https://github.com/MicrosoftDocs/azure-databases-docs/blob/main/articles/cosmos-db/analytics-and-business-intelligence-overview.md)
[databases-docs/blob/main/articles/cosmos-](https://github.com/MicrosoftDocs/azure-databases-docs/blob/main/articles/cosmos-db/analytics-and-business-intelligence-overview.md)
[db/analytics-and-business-intelligence-](https://github.com/MicrosoftDocs/azure-databases-docs/blob/main/articles/cosmos-db/analytics-and-business-intelligence-overview.md)
[overview.md](https://github.com/MicrosoftDocs/azure-databases-docs/blob/main/articles/cosmos-db/analytics-and-business-intelligence-overview.md). [Accessed 29 3 2025].